



DEV ∞ OPS PROJECTS — THAT — **ACTUALLY** LOOK LIKE PRODUCTION



- ✓ REAL-WORLD ARCHITECTURE
- ✓ INFRASTRUCTURE AS CODE
- ✓ CI/CD PIPELINE
- ✓ MONITORING & ALERTING
- ✓ SECURITY & RELIABILITY
- ✓ INCIDENTS & POSTMORTEMS



veriqta



veriqta



veriqta



veriqta



veriqta



veriqta

1. BUILD THE PRODUCTION ARCHITECTURE



DEFINE THE APPLICATION AND BUSINESS REQUIREMENTS

- Understand goals, features, performance and capacity needs
- Define SLAs, RTO, RPO and compliance requirements



USERS, FRONTEND, BACKEND, DATABASE

- Map user types and traffic
- Identify frontend clients, backend services and databases



PUBLIC VS PRIVATE COMPONENTS

- Identify internet-facing resources
- Keep sensitive components in private subnets



DNS AND ENTRY POINT

- Use Route 53 for DNS
- Define domain, records and routing policy



LOAD BALANCER

- Distribute traffic across healthy application targets
- Enable SSL/TLS



APPLICATION COMPUTE LAYER

- EC2 / Containers / ECS / EKS
- Auto Scaling for high availability
- Stateless application design



DATABASE LAYER

- Use managed DB
- Multi-AZ for high availability
- Automated backups



OBJECT STORAGE

- Store files, media, logs and backups
- Use S3 with proper encryption



CACHE

- Use ElastiCache (Redis / Memcached) for high performance and session storage



NETWORK BOUNDARIES

- VPC, Subnets, Route Tables
- Security Groups & NACLs
- Least privilege access



AVAILABILITY ZONES

- Distribute resources across multiple AZs for resilience



FAILURE DOMAINS

- AZs, subnets, instances
- Ensure no single point of failure



DRAW THE COMPLETE ARCHITECTURE BEFORE BUILDING


- Plan all components and their connections
- Review with stakeholders



IDENTIFY EVERY SINGLE POINT OF FAILURE

- List all dependencies
- Plan for redundancy
- Add monitoring and alerts

PROJECT ARCHITECTURE DIAGRAM (EXAMPLE)



REMEMBER: A strong architecture is planned before you build.
Good design prevents outages, reduces cost and simplifies operations.



YouTube



GitHub



LinkedIn



X



Instagram



Telegram

| veriqa

2. BUILD THE AWS NETWORK FOUNDATION

DESIGN A SECURE, RESILIENT AND SCALABLE PRODUCTION NETWORK



GOAL

Build a highly available VPC network that isolates resources, secures traffic and supports production workloads.

- ✓ PRODUCTION VPC
- ✓ CIDR PLANNING
- ✓ MULTIPLE AZs
- ✓ PUBLIC SUBNETS
- ✓ PRIVATE APP SUBNETS
- ✓ PRIVATE DB SUBNETS

- ✓ INTERNET GATEWAY
- ✓ NAT GATEWAY
- ✓ ROUTE TABLES
- ✓ SECURITY GROUPS
- ✓ NETWORK ACL AWARENESS
- ✓ PUBLIC vs PRIVATE IP



BEST PRACTICE

Keep private resources private.
Use least privilege security rules.
Test connectivity between all layers.



TEST CONNECTIVITY BETWEEN LAYERS


- EC2 in public subnet → Internet
- EC2 in private app subnet → Internet (via NAT)
- App subnet → DB subnet
- DB subnet → Internet (should FAIL)
- Use ping, curl, telnet and security group checks



TROUBLESHOOT A BROKEN ROUTE

- Check route table of the affected subnet
- Verify target (IGW, NAT, VGW, VPC Endpoint)
- Confirm next hop is correct
- Check NACLs and Security Groups
- Validate source/destination IPs and ports

PRODUCTION NETWORK DIAGRAM



ROUTE TABLES

- Public RT**
- 0.0.0.0/0 → IGW
- Private App RT**
- 0.0.0.0/0 → NAT GW
- Private DB RT**
- Local only



SECURITY GROUPS

- Allow only required traffic
- Restrict by source IP or SG
- No open to the world



NETWORK ACLs

- Stateless
- Subnet level
- Allow / Deny rules
- Use cautiously



CODE

Commit code



BUILD

Build container image



SCAN

Scan image and infrastructure



DEPLOY

Deploy to production



KEY TAKEAWAY

A strong network foundation is the backbone of secure, highly available AWS production workloads.

3. PROVISION EVERYTHING WITH TERRAFORM

Define. Provision. Version. Repeat.

TERRAFORM PROJECT STRUCTURE

```
terraform-project/
├── providers.tf
├── variables.tf
├── locals.tf
├── outputs.tf
├── modules/
│   ├── vpc/
│   ├── compute/
│   └── database/
├── environments/
│   ├── vpc/
│   ├── staging/
│   └── prod/
├── terraform.tfvars
└── backend.tf
```



PROVIDERS

Define the cloud or platform provider (e.g., AWS, Azure, GCP).



VARIABLES

Make configurations reusable and environment-aware.



LOCALS

Store expressions or values used across the configuration.



OUTPUTS

Expose values after provisioning (e.g., endpoints, IDs, ARNs).



REUSABLE MODULES

Organize infrastructure into reusable building blocks.

REMOTE STATE



Store Terraform state remotely for consistency and team collaboration.

S3 STATE STORAGE



Use an S3 bucket to store the state file securely.

STATE LOCKING



Prevent concurrent updates using DynamoDB locking.

ENVIRONMENT SEPARATION



Separate workspaces or directories for isolated environments.

VPC MODULE



Create VPC, subnets, route tables, IGW, NAT gateways, etc.

COMPUTE MODULE



Provision EC2, ASG, Launch Templates, Security Groups, IAM roles, etc.

DATABASE MODULE



Provision RDS, subnets, security groups, backups, parameter groups, etc.

TERRAFORM CLI WORKFLOW



terraform fmt
Format code for consistency.



terraform validate
Check configuration syntax and validity.



terraform plan
Preview changes before applying.



terraform apply
Apply the changes and provision infrastructure.

INFRASTRUCTURE CHANGE WORKFLOW



1. WRITE CODE
Define infrastructure as code.



2. PLAN
Review planned changes.



3. APPLY
Provision or update infrastructure.



4. VERIFY
Validate resources and outputs.

AVOID MANUAL INFRASTRUCTURE DRIFT



- ✓ No manual changes in console
- ✓ Everything defined in code
- ✓ Version controlled
- ✓ Repeatable and auditable
- ✓ Safe, consistent and scalable



Terraform ensures your infrastructure is consistent, repeatable and version controlled. Code is the source of truth.



YouTube



GitHub



LinkedIn



x



Instagram



Telegram



veriqta

4. BUILD A PRODUCTION GIT WORKFLOW

1. REPOSITORIES



APPLICATION REPOSITORY

Contains your application code.



INFRASTRUCTURE REPOSITORY

Contains IaC and environment configurations.

2. BRANCH PROTECTION



- ✓ Protect main, staging, production branches
- ✓ Require pull requests
- ✓ Require status checks to pass
- ✓ Require code review
- ✓ Prevent force push
- ✓ Prevent branch deletion

3. FEATURE BRANCHES



Create a branch from main for every new feature or fix.

4. PULL REQUESTS



Open a pull request to merge your changes into the target branch.

5. CODE REVIEW



At least one qualified reviewer must approve the changes.

6. REQUIRED STATUS CHECKS



All status checks must pass before merging (tests, linting, security scans, build).

7. CODEOWNERS



Define owners for paths and require their review on changes.

8. COMMIT DISCIPLINE



- ✓ Use meaningful commit messages
- ✓ Make small, focused commits
- ✓ One logical change per commit
- ✓ Follow team's commit message format

9. RELEASE TAGS



Tag every release in the repository (e.g., v1.2.0).

10. SEMANTIC VERSIONING

MAJOR	MINOR	PATCH
1	2	3

- ✓ MAJOR: Breaking changes
- ✓ MINOR: New features
- ✓ PATCH: Bug fixes



11. PREVENT DIRECT PRODUCTION CHANGES

Never push directly to production or main. Always go through pull requests.

12. HANDLE A REJECTED PULL REQUEST



- ✓ Read reviewer feedback carefully
- ✓ Update your branch
- ✓ Address all comments
- ✓ Request review again
- ✓ Never ignore feedback

13. ROLL BACK A BAD CHANGE



- ✓ Identify the last good release/tag
- ✓ Revert the commit or merge
- ✓ Deploy the previous good state
- ✓ Document what happened

14. GIT AS THE SOURCE OF CHANGE HISTORY



- ✓ Every change is versioned
- ✓ Full audit trail of who, what, when, why
- ✓ Essential for debugging, compliance and rollbacks



KEY TAKEAWAY

A strong Git workflow protects your production environment, ensures quality, enables safe collaboration and gives you complete control over every change.



CODE



REVIEW



MERGE



TAG



DEPLOY

5. CONTAINERIZE THE APPLICATION PROPERLY

Build secure, efficient and production-ready container images.

1. PRODUCTION DOCKERFILE

```
# Multi-stage build example
FROM node:20-alpine AS builder
WORKDIR /app
COPY package*.json ./
RUN npm ci
COPY . .
RUN npm run build

FROM node:20-alpine AS production
WORKDIR /app
ENV NODE_ENV=production
COPY --from=builder /app/dist ./dist
COPY package*.json ./
RUN npm ci --omit=dev
USER node
EXPOSE 3000
HEALTHCHECK --interval=30s --timeout=3s \
  CMD wget -qO- http://localhost:3000/health || exit 1
CMD ["node", "dist/server.js"]
```



2. MULTI-STAGE BUILDS



Build in one stage, run in another. Keep only what you need.

3. SMALL BASE IMAGES



Use minimal images like Alpine, Distroless or Slim.

4. DEPENDENCY MANAGEMENT



Install only production dependencies. Avoid bloat.

5. NON-ROOT CONTAINERS



Run as non-root user for better security.

6. .DOCKERIGNORE



Exclude files and folders you don't need in the image build.

7. ENVIRONMENT CONFIGURATION



Use ENV vars. Never hardcode secrets or configs.

8. HEALTH CHECKS



Define health checks so orchestrators know when your app is healthy.

9. IMAGE TAGGING



Use meaningful tags. Example: v1.2.0, 2.3.1 not latest.

10. IMMUTABLE IMAGES



Never change an image. Build a new image for every change.

11. VULNERABILITY SCANNING



Scan images for vulnerabilities before pushing to production.

12. BUILD AND INSPECT THE IMAGE

BUILD THE IMAGE



```
docker build -t myapp:1.0 .
```

INSPECT THE IMAGE



```
docker images myapp:1.0
docker inspect myapp:1.0
```

ANALYZE LAYERS AND SIZE



```
docker history myapp:1.0
docker image ls -s
```



13. RUN LOCALLY

```
docker run -d --name myapp \
-p 3000:3000 myapp:1.0
```

```
docker ps
docker logs myapp
```



14. DEBUG A CONTAINER THAT IMMEDIATELY EXITS

1. Check logs

```
docker logs myapp
```

2. Run interactively

```
docker run -it --rm myapp:1.0 sh
```

3. Check entrypoint, command, env vars and ports.



15. REDUCE UNNECESSARY IMAGE SIZE

- ✓ Use minimal base images
- ✓ Use multi-stage builds
- ✓ Remove dev dependencies
- ✓ Clean package caches
- ✓ Exclude files with .dockerignore
- ✓ Verify image size before pushing

THE RIGHT WORKFLOW



SCAN



FIX



REBUILD



RESCAN




WHY REBUILDING MATTERS

Changes in code or dependencies are not reflected in the old image.

Rebuild ensures every fix is actually included in the final image.

6. BUILD A CONTAINER REGISTRY WORKFLOW

Build, store and deliver container images securely and reliably.



7. VERSIONED IMAGE TAGS



- ✓ Use versioned tags for traceability and rollback.
- ✓ Examples:
 - v1.0.0
 - v1.2.3
 - 1.2.3
 - 2024-05-26
 - build-123

8. COMMIT SHA TAGGING

abc1234

Always tag images with the commit SHA for exact traceability.
Example: **abc1234**

9. AVOID RELYING ON LATEST



latest is mutable and can break deployments.
Always use specific tags.

10. IMAGE RETENTION POLICIES



- ✓ Keep only necessary versions.
- ✓ Delete untagged images.
- ✓ Use lifecycle policies (ECR lifecycle rules).
- ✓ Reduce storage costs.
- ✓ Retain images needed for rollback and compliance.

11. PULL IMAGE INTO ANOTHER ENVIRONMENT



- ✓ Authenticate with the registry.
- ✓ Pull the specific image tag.
- ✓ Verify image integrity.
- ✓ Deploy to target environment.

12. HANDLE AUTHENTICATION FAILURE



- ✗ Check credentials and permissions.
- ✗ Verify token expiration.
- ✗ Ensure correct region and registry URL.
- ✗ Re-authenticate and retry.

13. HANDLE MISSING IMAGE TAGS



- ✗ Verify the tag name.
- ✗ Check case sensitivity.
- ✗ Confirm tag was pushed.
- ✗ List available tags in registry.

14. TRACE DEPLOYMENT BACK TO SOURCE COMMIT



- ✓ Use commit SHA tagging.
- ✓ Record image tag in deployment manifests.
- ✓ Link deployment → image → commit.
- ✓ Enables debugging and audits.

WORKFLOW SUMMARY




A STRONG REGISTRY WORKFLOW = SECURE, TRACEABLE, AND RELIABLE DEPLOYMENTS

7. BUILD CI THAT DEVELOPERS CAN TRUST

FAST • RELIABLE • REPRODUCIBLE • SECURE

1. CI PIPELINE FLOW (HIGH LEVEL)



2. REQUIRED CHECKS BEFORE MERGE



- ✓ All jobs completed successfully
- ✓ No linting errors
- ✓ All unit tests passed
- ✓ All integration tests passed
- ✓ Build validation passed
- ✓ Security scans passed
- ✓ No critical or high vulnerabilities
- ✓ Artifacts generated
- ✓ Code review approved
- ✓ Up-to-date with target branch

3. PIPELINE PERFORMANCE BOOSTERS



PIPELINE CACHING

- Cache dependencies
- Cache Docker layers
- Cache build outputs
- Speeds up repeat runs



PARALLEL JOBS

- Run independent jobs in parallel
- Tests, lint, scan
- Reduce total time



OPTIMIZED AGENTS

- Right size runners
- Use faster machines
- Keep tools updated

4. READ PIPELINE LOGS SYSTEMATICALLY

- 1 OPEN THE FAILED JOB
Check the job that failed first.
- 2 READ FROM TOP TO BOTTOM
Start from the beginning of the logs.
- 3 FIND THE FIRST ERROR
Scroll up to the first actual error.
- 4 READ THE ERROR MESSAGE CAREFULLY
Understand what failed and why.
- 5 CHECK THE CONTEXT
Look at commands, inputs and environment.
- 6 IDENTIFY THE ROOT CAUSE
Focus on the real cause, not the last line.
- 7 TEST THE FIX LOCALLY
Reproduce locally and confirm the fix.
- 8 RE-RUN AND VERIFY
Push the fix and confirm the pipeline passes.



5. DIAGNOSE TEST THAT PASSES LOCALLY BUT FAILS IN CI



- ✓ CHECK ENVIRONMENT DIFFERENCES
OS, Node/Python version, databases, services, env variables.
- ✓ CHECK DEPENDENCIES
Different versions can cause failures.
- ✓ CHECK TEST DATA
Data setup may differ in CI.
- ✓ CHECK TIMING AND PARALLELISM
Tests that depend on order may fail in CI.
- ✓ CHECK EXTERNAL SERVICES
APIs, databases or third-party services.
- ✓ REPRODUCE CI ENVIRONMENT LOCALLY
Use Docker or the same CI image.
- ✓ ADD LOGS AND DEBUG OUTPUT
Print variables, states and responses.
- ✓ FIX THE ROOT CAUSE, NOT THE SYMPTOM
Avoid hardcoded values or workarounds.



A GOOD CI PIPELINE GIVES DEVELOPERS CONFIDENCE.
EVERY COMMIT, EVERY TIME.



YouTube



GitHub



LinkedIn



X



Instagram



Telegram



veriqta

8. BUILD THE CONTINUOUS DELIVERY PIPELINE

ONE BUILD. ONE ARTIFACT. MANY ENVIRONMENTS.



PIPELINE PRINCIPLES



RELIABILITY
Automate checks at every stage.



SECURITY
Scan and validate before promotion.



CONSISTENCY
The same artifact moves through all environments.



SPEED
Fast feedback, fast recovery, zero guesswork.

NEVER REBUILD AN ARTIFACT BETWEEN ENVIRONMENTS.



KEY TAKEAWAY


Build once, test everywhere, promote with confidence, and deliver value to users safely and consistently.



9. DEPLOY THE APPLICATION TO KUBERNETES

Deploy, manage and operate production-ready applications on Kubernetes.

KUBERNETES ARCHITECTURE OVERVIEW



CORE K8S OBJECTS

- NAMESPACE**
Logical isolation for resources.
- DEPLOYMENT**
Manages replica sets and updates.
- PODS**
Smallest deployable units that run containers.
- SERVICE**
Stable network endpoint to access pods.
- CONFIGMAPS**
Store non-sensitive configuration data.
- SECRETS**
Store sensitive data such as passwords, tokens and keys.

RESOURCE MANAGEMENT

RESOURCE REQUESTS



Minimum resources the container needs.

```
resources:
  requests:
    cpu: "250m"
    memory: "256Mi"
```

RESOURCE LIMITS



Maximum resources the container can use.

```
resources:
  limits:
    cpu: "500m"
    memory: "512Mi"
```

WHY IT MATTERS



- ✓ Prevents noisy neighbors
- ✓ Better scheduling
- ✓ Avoids OOMKills
- ✓ Improves stability

HEALTH CHECKS (PROBES)



LIVENESS PROBE

Checks if the container is alive. Restarts the container if it fails.

```
livenessProbe:
  httpGet:
    path: /healthz
    port: 8080
```



READINESS PROBE

Checks if the container is ready to receive traffic.

```
readinessProbe:
  httpGet:
    path: /ready
    port: 8080
```



STARTUP PROBE

Gives slow-starting containers time to initialize.

```
startupProbe:
  httpGet:
    path: /startup
    port: 8080
```

UPDATES & AVAILABILITY



ROLLING UPDATES

Update pods gradually without downtime.



REPLICA STRATEGY

Maintain the desired number of pod replicas.



POD DISRUPTION AWARENESS

Distribute pods across nodes and zones to maintain availability.

DEPLOY & INSPECT WORKLOADS

- Apply: `kubectl apply -f deployment.yaml`
- Get Pods: `kubectl get pods -n <namespace>`
- Describe: `kubectl describe pod <pod-name>`
- Logs: `kubectl logs <pod-name>`
- Exec: `kubectl exec -it <pod-name> -- /bin/sh`

DEBUG COMMON FAILURES



CrashLoopBackOff

- Container crashes repeatedly.
- Check logs: `kubectl logs <pod-name>`
 - Fix the application or configuration.



ImagePullBackOff

- Kubernetes cannot pull the image.
- Check image name and registry.
 - Verify image pull secrets and access.



KEY TAKEAWAY

Deploy with best practices, monitor your workloads, and debug issues fast to keep applications reliable in production.



10. EXPOSE THE APPLICATION SAFELY

KEY CONCEPTS



INTERNAL vs EXTERNAL SERVICES

Internal services are only accessible inside the cluster. External services are reachable from outside.



INGRESS

Kubernetes resource that defines rules for external access to services.



INGRESS CONTROLLER

Software that watches Ingress resources and configures the proxy/load balancer.



CLOUD LOAD BALANCER

Distributes traffic across nodes running the Ingress Controller (usually a Service of type LoadBalancer).



DNS

Maps your domain name to the public IP or hostname of the load balancer.



TLS / HTTPS

Encrypts traffic in transit to protect users and data.



CERTIFICATE MANAGEMENT

Use cert-manager or cloud provider tools to issue and renew TLS certificates.



HOST-BASED ROUTING

Route traffic based on the domain name (e.g. app.example.com).



PATH-BASED ROUTING

Route traffic based on URL path (e.g. /api, /web).



SECURITY GROUPS

Control inbound and outbound traffic to allow only what is necessary.

END-TO-END REQUEST FLOW



USER



DNS

Resolves domain to Load Balancer IP



LOAD BALANCER

Receives external traffic and forwards to Ingress Controller



INGRESS

Applies rules and routes traffic to the correct Service



SERVICE

Routes traffic to the appropriate Pods (endpoints)



POD

Application processes the request and returns response

RESPONSE RETURNS THE SAME PATH



DEBUG A 502 RESPONSE

- CHECK POD HEALTH**
Ensure Pods are Running and Ready.
`kubectl get pods -n <namespace>`
- CHECK SERVICE ENDPOINTS**
Ensure the Service has healthy endpoints.
`kubectl get endpoints -n <namespace> <service>`
- CHECK INGRESS CONTROLLER LOGS**
Look for upstream errors or connection issues.
`kubectl logs -n <ingress-namespace> <ingress-pod>`
- CHECK BACKEND POD LOGS**
Verify the application is not crashing or returning errors.
- CHECK TIMEOUT SETTINGS**
Increase timeouts if your application takes longer to respond.
- CHECK SECURITY GROUPS / NETWORK**
Ensure traffic from Load Balancer to Nodes is allowed.



APP WORKS INSIDE CLUSTER BUT NOT EXTERNALLY

- CHECK DNS RESOLUTION**
Does the domain resolve to the correct Load Balancer IP?
- VERIFY LOAD BALANCER HEALTH**
Is the Load Balancer healthy and targeting the nodes?
- CHECK INGRESS RULES**
Host/path rules may be incorrect or missing.
- VERIFY SERVICE TYPE**
Ensure the Service is ClusterIP (for Ingress) and exposed correctly.
- CHECK SECURITY GROUPS / FIREWALL**
Ensure inbound 80/443 allowed to Load Balancer and NodePort (if used).
- CHECK TLS / CERTIFICATE**
Certificate mismatch or expired cert can break HTTPS.
- CHECK PROXY / APPLICATION CONFIG**
Ensure the app is aware of correct host/protocol headers.

KEY TAKEAWAYS



- Expose only what is necessary.
- Use TLS everywhere.
- Use Ingress for flexible routing.
- Keep Security Groups tight.
- Monitor and check logs.
- Test from outside like a real user.
- Follow the end-to-end path when debugging.
- A 502 usually means upstream or connectivity issues.



YouTube



GitHub



LinkedIn



Instagram



Telegram

| veriQTA



11. ADD PRODUCTION CONFIGURATION AND SECRETS

CONFIGURATION vs SECRETS



CONFIGURATION

- Non-sensitive settings like URLs, feature flags, timeouts, log levels.
- Can be public.

VS



SECRETS

- Sensitive data like passwords, API keys, tokens, certificates.
- Must be protected, encrypted and access controlled.

ENVIRONMENT-SPECIFIC CONFIGURATION



- ✓ Different values for Dev, Staging and Production.
- ✓ Use environment variables or configuration systems.
- ✓ Keep code the same, change configuration.

KUBERNETES CONFIGMAPS



- Store non-sensitive configuration.
- Key-value pairs.
- Mounted as files or environment variables.

KUBERNETES SECRETS



- Store sensitive data (Base64 encoded).
- Used as environment variables or files.
- Access controlled by RBAC.

EXTERNAL SECRET STORES



Use managed secret stores for better security, auditing and rotation.

AWS SECRET STORES



AWS SECRETS MANAGER

- Store, rotate and manage secrets.
- Automatic rotation support.
- Fine-grained IAM access.



AWS PARAMETER STORE

- Store configuration and secrets as SecureString.
- Hierarchical naming.
- Lower cost for simple needs.

SECRET INJECTION



- Inject secrets into Pods at runtime.
- Use environment variables or mounted files.
- Avoid baking secrets into images.



IAM-BASED ACCESS

- Grant least privilege IAM permissions.
- Workload identity for Pods and services.
- Audit access with CloudTrail.



SECRET ROTATION

- Rotate secrets regularly.
- Reduce risk if a secret is exposed.
- Test new secret before revoking the old one.



AVOID CREDENTIALS IN GIT

- Never commit secrets or keys.
- Use .gitignore.
- Scan for secrets before commit.



AVOID SECRETS IN DOCKER IMAGES

- Do not store secrets in Dockerfile or image layers.
- Use runtime injection instead.
- Images are shared and cached.



HANDLE MISSING CONFIGURATION

- Application should fail fast.
- Clear error messages.
- Use defaults only for non-sensitive values.



HANDLE AN EXPIRED CREDENTIAL

- Detect expiration early.
- Implement automatic refresh or rotation.
- Alert and retry securely.

REMOVE AN ACCIDENTALLY COMMITTED SECRET SAFELY



- 1 Rotate or revoke the exposed secret immediately.
- 2 Remove the secret from code history (BFG Repo-Cleaner or filter-repo).
- 3 Force push the cleaned history.
- 4 Invalidate caches and CI/CD artifacts.
- 5 Notify your team and audit access logs.



KEY TAKEAWAY: PROTECT CONFIGURATION AND SECRETS EVERYWHERE. STORE SAFELY, ACCESS SECURELY, ROTATE REGULARLY.



YouTube



GitHub



LinkedIn



X



Instagram



Telegram


| veriqta

12. BUILD THE DATABASE LAYER



1. DATABASE ARCHITECTURE

- ✓ Managed relational database (AWS RDS / Aurora / Cloud SQL)
- ✓ Private database subnet
- ✓ Application-to-database connectivity
- ✓ Multi-AZ deployment for high availability



2. SECURITY

- ✓ Database security groups (allow only required ports from application layer)
- ✓ Strong credentials (no hardcoded secrets)
- ✓ Use Secrets Manager / Parameter Store
- ✓ Encrypt in transit (TLS) and at rest



3. CONNECTION MANAGEMENT

- ✓ Use secure connection strings
- ✓ Enable connection pooling
- ✓ Set max connections appropriately
- ✓ Monitor connection usage
- ✓ Prevent connection exhaustion



4. SCHEMA & MIGRATIONS

- ✓ Use version-controlled migrations
- ✓ Automate schema migrations
- ✓ Test migrations in staging
- ✓ Avoid manual schema changes in production



5. BACKUPS & RECOVERY

- ✓ Enable automated backups
- ✓ Configure backup retention
- ✓ Use point-in-time recovery (PITR)
- ✓ Test restore regularly
- ✓ Store backups in secure location



6. HIGH AVAILABILITY

- ✓ Multi-AZ deployment enabled
- ✓ Automatic failover configured
- ✓ Monitor replication lag
- ✓ Test failover behavior periodically



7. ENCRYPTION

- ✓ Encryption at rest (KMS)
- ✓ Encryption in transit (TLS)
- ✓ Manage keys securely
- ✓ Rotate keys when required



8. MONITORING

- ✓ Enable database monitoring
- ✓ Track CPU, memory, storage, IOPS
- ✓ Set up alerts for key metrics
- ✓ Integrate with CloudWatch / Monitoring tools



9. PERFORMANCE

- ✓ Monitor slow queries
- ✓ Use indexes effectively
- ✓ Analyze query performance
- ✓ Optimize regularly



10. COMMON ISSUES

- ✓ Slow queries
- ✓ Connection exhaustion
- ✓ High CPU / memory
- ✓ Lock contention
- ✓ Replication lag



11. RESILIENCE & INCIDENT HANDLING

- ✓ Simulate database unavailability
- ✓ Failover to standby (Multi-AZ)
- ✓ Restore from backup if needed
- ✓ Verify application connectivity
- ✓ Recover without blindly restarting everything
- ✓ Communicate and validate data integrity



KEY TAKEAWAY: A secure, highly available and well-monitored database layer is critical for application reliability and data durability.



YouTube



GitHub



LinkedIn



Instagram



Telegram

| veriqta


13. ADD CACHING AND ASYNCHRONOUS PROCESSING

REDIS CACHE



- Fast in-memory key-value store
- Great for reducing database load and improving response time
- Used for sessions, tokens, metadata, query results, etc.

CACHE-ASIDE PATTERN



CACHE HIT vs MISS



CACHE HIT

Data found in Redis.



CACHE MISS

Data not in cache. Fetch from DB, then store in Redis.

TTL (TIME TO LIVE)



Each key expires after a set time to keep data fresh.

CACHE INVALIDATION



Remove or expire stale data when underlying data changes.

BEST PRACTICES


- ✓ Use appropriate TTL
- ✓ Invalidate on updates
- ✓ Avoid caching sensitive data
- ✓ Monitor hit ratio

MESSAGE QUEUE



Enables asynchronous processing and decouples services.

HOW IT WORKS



PRODUCERS



Send messages to the queue.

CONSUMERS



Read messages from the queue.

BACKGROUND WORKERS



Process messages asynchronously.

RETRIES



Retry failed messages automatically.

DEAD-LETTER QUEUE (DLQ)



Stores messages that continuously fail after max retries.

IDEMPOTENCY



Ensure operations can be processed multiple times safely.

QUEUE DEPTH



Number of messages waiting. High depth may indicate backlog.


POISON MESSAGES




Bad messages that always fail. Send to DLQ to prevent blocking the queue.

ARCHITECTURE OVERVIEW

APPLICATION → REDIS



APPLICATION → QUEUE → WORKER



TROUBLESHOOT: GROWING QUEUE BACKLOG

- Check worker health and number of active workers.
- Check message processing time.
- Check for slow downstream services or DB.
- Increase workers or auto-scaling.
- Check for poison messages (look at DLQ).
- Monitor queue depth and set alerts.



WHAT HAPPENS WHEN REDIS DISAPPEARS?

- All cache lookups fail.
- Applications fall back to DB (cache miss).
- Increased DB load and higher response time.
- Sessions may be lost (if stored in cache).
- Temporary failures until Redis returns.
- Design apps to be cache-resilient (graceful fallback).



KEY TAKEAWAY

Caching makes things fast. Queues make things reliable. Use both together to build scalable, resilient applications.



YouTube



GitHub



LinkedIn



X



Instagram



Telegram



veriqta

14. BUILD FULL OBSERVABILITY

 **OBSERVE EVERYTHING. UNDERSTAND FAST. FIX WITH CONFIDENCE.**



METRICS

Numbers that show the health and performance.

- ✓ Application metrics
- ✓ Infrastructure metrics
- ✓ Kubernetes metrics
- ✓ Request latency
- ✓ Error rate
- ✓ Traffic
- ✓ Saturation



LOGS

Detailed events that explain what happened.

- ✓ Centralized logging
- ✓ Structured logs
- ✓ Log aggregation
- ✓ Search and filter
- ✓ Retain and archive
- ✓ Alert on important logs



TRACES

Follow a request as it moves across services.

- ✓ Distributed tracing
- ✓ Service map
- ✓ Span timeline
- ✓ Latency breakdown
- ✓ Bottleneck detection
- ✓ Root cause analysis

KEY OBSERVABILITY TOOLS



PROMETHEUS

Metrics collection and alerting.



GRAFANA

Dashboards and visualization.



CLOUDWATCH

AWS native metrics, logs and alerts.



KUBERNETES METRICS

Cluster, node, pod and workload metrics.

OBSERVABILITY ESSENTIALS



CORRELATION IDS

Track a single request across all services.



USEFUL DASHBOARDS

Show what matters. Clear, simple and actionable.



ACTIONABLE ALERTS

Alert on symptoms that indicate real problems.



REDUCE NOISE

Focus on high signal alerts, not too many false alarms.



BUSINESS CONTEXT


Understand impact on users, revenue and systems.



SPEED TO INSIGHT

Find the issue fast. Fix faster.

FOLLOW ONE FAILED REQUEST ACROSS THE SYSTEM



TRACES + LOGS + METRICS + CORRELATION ID = ROOT CAUSE FOUND



**FULL OBSERVABILITY = SEE EVERYTHING, UNDERSTAND EVERYTHING, FIX FASTER.
BETTER INSIGHT. BETTER RELIABILITY. BETTER SYSTEMS.**



YouTube



GitHub



LinkedIn



X



Instagram



Telegram

| veriqta

15. ENGINEER RELIABILITY AND SELF-HEALING



**BUILD SYSTEMS THAT HEAL THEMSELVES,
NOT SYSTEMS THAT BREAK AND ALERT.**

CORE RELIABILITY PATTERNS



HORIZONTAL POD AUTOSCALER

Automatically scales pods based on CPU, memory or custom metrics.



REPLICA STRATEGY

Run multiple replicas of your application to ensure high availability.



POD ANTI-AFFINITY

Spread pods across different nodes to avoid single node failure.



TOPOLOGY SPREAD

Distribute pods evenly across zones, nodes or other topology domains.



MULTIPLE AVAILABILITY ZONES

Run workloads across multiple AZs to survive a zone failure.



POD DISRUPTION BUDGETS

Limit the number of pods that can be unavailable during disruptions.



HEALTH PROBES

Liveness, readiness and startup probes ensure only healthy pods receive traffic.



GRACEFUL TERMINATION

Give applications time to finish work before a pod is stopped.



ROLLING UPDATES

Update pods gradually to prevent downtime and allow easy rollback.



RETRY BEHAVIOR

Retry failed requests with backoff to handle temporary failures.



TIMEOUTS

Set timeouts to prevent hanging requests from consuming resources.



CIRCUIT BREAKERS

Stop sending requests to failing dependencies to protect your system.



DEPENDENCY FAILURES

Design for failures. Assume dependencies will fail.

CHAOS TEST YOUR SYSTEM



REMOVE ONE POD INTENTIONALLY

Observe how the system recreates and restores it.



REMOVE ONE NODE INTENTIONALLY

Observe pod rescheduling and traffic continuity.

SELF-HEALING WORKFLOW



DETECT

Health probes, metrics and alerts detect an issue.



ISOLATE

Kubernetes removes unhealthy pods from service.



REPLACE

New pods are scheduled on healthy nodes.



RECOVER

Traffic resumes automatically. System stays available.



LEARN & IMPROVE

Review, improve configuration and eliminate weak points.



IDENTIFY REMAINING SINGLE POINTS OF FAILURE

- Single database instance
- Stateful workloads without backups
- External APIs without fallbacks
- No multi-AZ deployment
- Critical dependencies without failover
- Shared storage or network bottlenecks



YouTube



GitHub



LinkedIn



X



Instagram



Telegram

| veriqta

16. SECURE THE DELIVERY AND RUNTIME ENVIRONMENT



SECURE BY DESIGN. SECURE BY DEFAULT. VERIFY ALWAYS.



LEAST PRIVILEGE

- Grant only the access required
- Remove unnecessary permissions
- Review regularly



IAM ROLES

- Use IAM roles, not long-term keys
- Scope roles narrowly
- Rotate and audit regularly



WORKLOAD IDENTITY

- Use OIDC / IRSA for Kubernetes
- No static credentials in workloads
- Short-lived tokens



CI/CD PERMISSIONS

- Minimal pipeline permissions
- Separate build and deploy roles
- Isolate environments



DEPENDENCY SCANNING

- Scan dependencies early and often
- Catch vulnerable libraries
- Keep dependencies up to date



SECRET SCANNING

- Detect secrets in code and configs
- Prevent secret leakage
- Rotate exposed secrets



CONTAINER IMAGE SCANNING

- Scan images for vulnerabilities
- Use minimal base images
- Rescan after updates



INFRASTRUCTURE SCANNING

- Scan IaC (Terraform, CloudFormation, K8s, etc.)
- Find misconfigurations
- Prevent insecure deployments



KUBERNETES SECURITY CONTEXT

- Set securityContext in all Pods
- Drop capabilities
- Read-only root filesystem



NON-ROOT CONTAINERS

- Run containers as non-root
- Set runAsUser
- Reduce impact of compromise



NETWORK POLICIES

- Control Pod-to-Pod traffic
- Allow only required communication
- Default deny is safer



TLS EVERYWHERE

- Encrypt data in transit
- Enforce TLS 1.2 or higher
- Use valid certificates



ENCRYPTION AT REST

- Encrypt databases, volumes, backups
- Use KMS
- Protect data at all layers



RBAC

- Use Role-Based Access Control
- Grant least privilege access
- Review bindings regularly



AUDITABILITY

- Enable audit logs
- Monitor activities
- Centralize logs
- Detect and respond to anomalies



SOFTWARE SUPPLY CHAIN AWARENESS

- Know your dependencies
- Verify sources and signatures
- Use trusted registries
- Monitor for new vulnerabilities



TEST YOUR CONTROLS

ATTEMPT AN UNAUTHORIZED OPERATION



VERIFY THAT SECURITY CONTROLS BLOCK IT

EXAMPLES

- ✓ Access a restricted resource
- ✓ Escalate privileges
- ✓ Exfiltrate secrets
- ✓ Deploy a privileged pod



SECURITY IS EVERYONE'S RESPONSIBILITY.
BUILD SECURE. DEPLOY SECURE. RUN SECURE.

17. BUILD ZERO-DOWNTIME DEPLOYMENT AND ROLLBACK

1. DEPLOYMENT STRATEGIES

ROLLING DEPLOYMENT



Replace old instances with new ones gradually.
No downtime.

BLUE-GREEN DEPLOYMENT



Run two identical environments.
Switch traffic to Green in one step.
Instant rollback to Blue.

CANARY DEPLOYMENT



Send a small % of traffic to the new version first. Increase gradually if healthy.

2. TRAFFIC AND HEALTH



TRAFFIC SHIFTING

Move traffic between versions using ALB, Ingress, or Service Mesh.



HEALTH VERIFICATION

Check readiness, liveness, and external health endpoints.



DEPLOYMENT METRICS

Monitor latency, throughput, CPU, memory, and requests.



ERROR-RATE COMPARISON

Compare error rates between new and old versions.

3. ROLLBACK AUTOMATION



AUTOMATED ROLLBACK CRITERIA

Rollback when error rate, latency, or failed health checks exceed set thresholds.



ROLLBACK APPLICATION VERSION

Revert to the previous stable version automatically.



ROLL-FORWARD STRATEGY

If rollback is not possible, quickly deploy a fixed version forward.

4. DATA AND COMPATIBILITY



DATABASE COMPATIBILITY

Ensure schema changes are backward and forward compatible.



BACKWARD-COMPATIBLE CHANGES

New releases must work with old data and old services.



FEATURE FLAGS

Hide new features behind flags. Enable gradually and safely.

5. PRACTICAL ZERO-DOWNTIME WORKFLOW



1. DEPLOY

Deploy a deliberately faulty release.



2. DETECT

Detect degradation using metrics and alerts.



3. ANALYZE

Confirm the issue (metrics, logs, traces).



4. DECIDE

Decide: Rollback or Roll-forward.



5. RESTORE

Restore service safely with zero or minimal impact.

BEST PRACTICES



- ✓ Always verify health before shifting more traffic.
- ✓ Automate rollbacks to reduce human delay.
- ✓ Use small increments in canary deployments.
- ✓ Keep deployments idempotent and repeatable.
- ✓ Test rollback process regularly.
- ✓ Monitor user experience, not just infrastructure.

ZERO-DOWNTIME CHECKLIST



- ✓ Choose deployment strategy (Rolling / Blue-Green / Canary)
- ✓ Verify health and metrics
- ✓ Check database and backward compatibility
- ✓ Enable feature flags
- ✓ Monitor error rate, latency, and throughput
- ✓ Have automated rollback in place



GOAL: Deliver new versions continuously while keeping users happy.
Detect problems fast. Recover automatically. Keep services available.



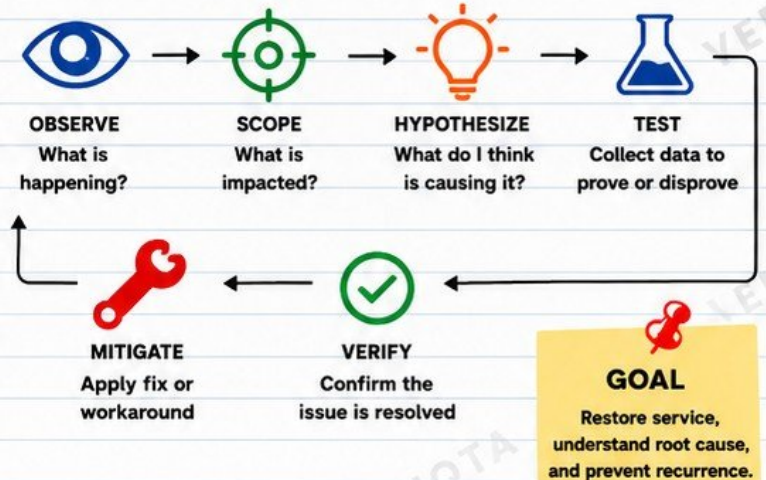
BREAK PRODUCTION ON PURPOSE

BUILD RESILIENCE. FIND WEAKNESSES. IMPROVE BEFORE USERS DO.

CONTROLLED FAILURE SCENARIOS

1. POD CRASH
2. NODE FAILURE
3. CPU SATURATION
4. MEMORY EXHAUSTION
5. DISK PRESSURE
6. DNS FAILURE
7. EXPIRED SECRET
8. DATABASE CONNECTION EXHAUSTION
9. BROKEN SECURITY GROUP
10. BAD KUBERNETES CONFIGURATION
11. FAILED DEPLOYMENT
12. QUEUE BACKLOG
13. DEPENDENCY TIMEOUT
14. HIGH APPLICATION LATENCY
15. HTTP 5XX SPIKE

THE INCIDENT INVESTIGATION FLOW



FOR EVERY INCIDENT, CAPTURE THIS

- ✓ TIMELINE (WHAT HAPPENED AND WHEN)
- ✓ IMPACT (WHO WAS AFFECTED)
- ✓ DETECTION METHOD
- ✓ OBSERVATIONS AND EVIDENCE
- ✓ COMMANDS RUN
- ✓ ROOT CAUSE
- ✓ MITIGATION TAKEN
- ✓ LESSONS LEARNED

USEFUL DIAGNOSTIC COMMANDS

WHAT TO CHECK	COMMAND EXAMPLE
POD STATUS	<code>kubectl get pods -A</code>
POD DETAILS	<code>kubectl describe pod <pod-name> -n <ns></code>
POD LOGS	<code>kubectl logs <pod-name> -n <ns></code>
NODE STATUS	<code>kubectl get nodes</code>
NODE ISSUES	<code>kubectl describe node <node-name></code>
EVENTS	<code>kubectl get events -A --sort-by=.metadata.creationTimestamp</code>
RESOURCE USAGE	<code>kubectl top pods -A</code>
STORAGE	<code>df -h</code>
NETWORK	<code>kubectl get svc,ing,ep -A</code>
DNS	<code>nslookup <service-name></code> <code>dig <domain></code>
SECRETS	<code>kubectl get secrets -A</code>
DEPLOYMENTS	<code>kubectl rollout status deployment/<name> -n <ns></code>

FAILURE LAB EXAMPLE

EXAMPLE: POD CRASH LOOP

1. **BREAK IT:** Update deployment with invalid config or wrong image.
2. **OBSERVE:** `kubectl get pods -n <ns>` (See `CrashLoopBackOff`)
3. **SCOPE:** Which service is affected? Check metrics and logs.
4. **HYPOTHEZIZE:** Wrong image? Bad config? Missing secret?
5. **TEST:** Check logs, describe pod, events.
6. **MITIGATE:** Rollback or fix configuration.
7. **VERIFY:** Pods running? Metrics normal? Users not impacted?



TIP:

Break things in lower environments first. Document everything as if it were real.

REMEMBER



- ✓ FAIL FAST IN SAFE ENVIRONMENTS.
- ✓ AUTOMATE VALIDATION.
- ✓ BUILD SELF-HEALING SYSTEMS.
- ✓ OBSERVABILITY FIRST.
- ✓ DOCUMENT YOUR RECOVERY.
- ✓ TURN INCIDENTS INTO IMPROVEMENTS.



REAL ENGINEERS DON'T FEAR FAILURES. THEY USE THEM TO BUILD RELIABLE SYSTEMS.



veriqta



veriqta



veriqta



veriqta



veriqta



veriqta

RUN A REAL INCIDENT AND WRITE THE POSTMORTEM

NO BLAME. EVIDENCE FIRST.

SCENARIO



A DEPLOYMENT SUCCEEDS.

FIVE MINUTES LATER:



HTTP 5xx INCREASES



P95 LATENCY RISES



SOME PODS RESTART



DATABASE CONNECTIONS INCREASE



CUSTOMERS REPORT INTERMITTENT FAILURES



**SOMETHING IS WRONG.
LET'S INVESTIGATE.**

INVESTIGATE



DEPLOYMENT HISTORY



KUBERNETES EVENTS



POD LOGS



METRICS



TRACES



DATABASE STATE



RECENT CONFIGURATION CHANGES



GATHER EVIDENCE BEFORE FORMING CONCLUSIONS.

START WITH OBSERVABILITY



METRICS

Check error rate, latency, traffic, CPU, memory, DB connections.



LOGS

Look for errors, exceptions, timeouts, restarts.



TRACES

Follow a request end-to-end.



DATABASE

Connections, slow queries, locks, resource usage.



CONFIG CHANGES

Recent deployments, feature flags, secrets, config updates.

INVESTIGATION FLOW (FOLLOW THIS LOOP)



1. OBSERVE

What is happening?



2. SCOPE

What is impacted?



3. HYPOTHEZIZE

What could be causing it?



4. TEST

Collect data to prove or disprove.



5. MITIGATE

Apply fix or workaround.



6. VERIFY

Confirm the issue is resolved.

POSTMORTEM: WHAT TO PRODUCE



INCIDENT TIMELINE

A clear timeline of what happened and when.



CUSTOMER IMPACT

Who was affected and how.



DETECTION METHOD

How was the incident detected?



ROOT CAUSE

The primary reason the incident happened.



CONTRIBUTING FACTORS

Other factors that made the incident worse.



MITIGATION

What actions reduced the impact?



RECOVERY

How was full service restored?



CORRECTIVE ACTIONS

What will we fix in the system?



PREVENTIVE ACTIONS

What will we do to avoid this in the future?



LESSONS LEARNED

What did we learn as a team?

EVIDENCE FIRST

- NO ASSUMPTIONS
- NO BLAME
- NO GUESSING
- JUST FACTS
- JUST IMPACT
- JUST IMPROVEMENT

PRO TIP

WRITE POSTMORTEMS WHILE IT'S FRESH. ACCURACY MATTERS. CLARITY HELPS EVERYONE IMPROVE.



THE GOAL IS NOT TO FIND SOMEONE TO BLAME.

THE GOAL IS TO IMPROVE THE SYSTEM SO IT DOESN'T HAPPEN AGAIN.



veriQTA



veriQTA



veriQTA



veriQTA



veriQTA




veriQTA


FINAL PRODUCTION PLATFORM PROJECT

BUILD. DELIVER. OPERATE. IMPROVE.

END-TO-END SYSTEM ARCHITECTURE




INFRASTRUCTURE (AS CODE)



DELIVERY PIPELINE



OPERATIONS & OBSERVABILITY



FINAL REQUIREMENTS (PRODUCTION READINESS CHECKLIST)

- | | | |
|--------------------------|----------------------------|------------------------------|
| ✓ INFRASTRUCTURE AS CODE | ✓ AUTOSCALING | ✓ FAILURE TESTING |
| ✓ MULTI-AZ ARCHITECTURE | ✓ MONITORING | ✓ INCIDENT RESPONSE |
| ✓ AUTOMATED CI/CD | ✓ ALERTING | ✓ ARCHITECTURE DOCUMENTATION |
| ✓ SECURE SECRET HANDLING | ✓ BACKUP AND RECOVERY | ✓ RUNBOOK |
| ✓ CONTAINER SECURITY | ✓ ZERO-DOWNTIME DEPLOYMENT | ✓ POSTMORTEM PROCESS |
| ✓ HIGH AVAILABILITY | ✓ ROLLBACK | ✓ FINAL PRODUCTION REVIEW |

PRODUCTION IS NOT A DESTINATION. IT IS A DISCIPLINE.

CONTINUOUSLY MEASURE.
CONTINUOUSLY IMPROVE.
CONTINUOUSLY PROTECT.



THE FINAL QUESTION IS NOT:

"DOES THE APPLICATION RUN?" →

THE REAL QUESTION IS:

"CAN ANOTHER ENGINEER SAFELY DEPLOY, OBSERVE, TROUBLESHOOT, RECOVER, AND OPERATE THIS SYSTEM WHEN THINGS GO WRONG?"

